

Lecture 23

# EECS 483: COMPILER CONSTRUCTION

# Announcements

- HW6: Analysis and Optimization
  - Due on Thursday, May 2
- Final Exam
  - 4-6pm April 29
  - DOW1010, DOW1005, DOW2166
  - Same cheat sheet policy as before
- Guest Lectures
  - Professor Cyrus Omar today
  - Steven Schaefer (GSI last semester) Wednesday
  - Professor New returns next week.




# LOOPS AND DOMINATORS

# Loops in Control-flow Graphs

- Taking into account loops is important for optimizations.
  - The 90/10 rule (Pareto Principle) applies, so optimizing loop bodies is important
- Should we apply loop optimizations at the AST level or at a lower representation?
  - Loop optimizations benefit from other IR-level optimizations and vice-versa, so it is good to interleave them.
- Loops may be hard to recognize at the quadruple / LLVM IR level.
  - Many kinds of loops: while, do/while, for, continue, goto...
- Problem: *How do we identify loops in the control-flow graph?*


# Definition of a (Structured) Loop

- A **structured loop** is a set of nodes in the control flow graph.
  - One distinguished entry point called the *header*
  - Every node is reachable from the header & the header is reachable from every node (using only nodes in the loop).
  - No edges enter the loop except to the header
- Nodes with outgoing edges are called loop exit nodes
- A loop is a *strongly connected component*




# Nested Loops

- Control-flow graphs may contain many loops
- Loops may contain other loops:




Control Tree:



The control tree depicts the nesting structure of the program.


# Control-flow Analysis

- Goal: Identify the loops and nesting structure of the CFG.
- Control flow analysis is based on the idea of *dominators*:
- Node A *dominates* node B if the only way to reach B from the start node is through node A.
- An edge in the graph is a *back edge* if the target node dominates the source node.
- A loop contains at least one back edge.



# Dominator Trees

- Domination is reflexive:
  - A dominates A
- Domination is transitive:
  - if A dominates B and B dominates C then A dominates C
- Domination is anti-symmetric:
  - if A dominates B and B dominates A then A = B
- Every flow graph has a dominator tree
  - The Hasse diagram of the dominates relation






# Dominator Dataflow Analysis

- We can define  $\text{Dom}[n]$  as a forward dataflow analysis.
  - Using the framework we saw earlier:  $\text{Dom}[n] = \text{out}[n]$  where:
- “A node B is dominated by another node A if A dominates *all* of the predecessors of B.”
  - $\text{in}[n] := \bigcap_{n' \in \text{pred}[n]} \text{out}[n']$
- “Every node dominates itself.”
  - $\text{out}[n] := \text{in}[n] \cup \{n\}$
- Formally:  $L = \text{set of nodes ordered by } \subseteq$ 
  - $T = \{\text{all nodes}\}$
  - $F_n(x) = x \cup \{n\}$
  - $\sqcap$  is  $\cap$
- Easy to show monotonicity and that  $F_n$  distributes over meet.
  - So algorithm terminates and is MOP (meet over all paths)



# Improving the Algorithm

- Dom[b] contains just those nodes along the path in the dominator tree from the root to b:
  - e.g. Dom[8] = {1,2,4,8}, Dom[7] = {1,2,4,5,7}
  - There is a lot of sharing among the nodes
- More efficient way to represent Dom sets is to store the dominator *tree*.
  - doms[b] = immediate dominator of b
  - doms[8] = 4, doms[7] = 5
- To compute Dom[b] walk through doms[b]
- Need to efficiently compute intersections of Dom[a] and Dom[b]
  - Traverse up tree, looking for least common ancestor:
  - Dom[8]  $\cap$  Dom[7] = Dom[4]
- See: “A Simple, Fast Dominance Algorithm” Cooper, Harvey, and Kennedy




# Completing Control-flow Analysis

- Dominator analysis identifies *back edges*:
  - Edge  $n \rightarrow h$  where  $h$  dominates  $n$
- Each back edge has a *natural loop*:
  - $h$  is the header
  - All nodes reachable from  $h$  that also reach  $n$  without going through  $h$
- For each back edge  $n \rightarrow h$ , find the natural loop:
  - $\{n' \mid n \text{ is reachable from } n' \text{ in } G - \{h\}\} \cup \{h\}$
- Two loops may share the same header: merge them
- Nesting structure of loops is determined by set inclusion
  - Can be used to build the control tree




# Example Natural Loops



Natural Loops



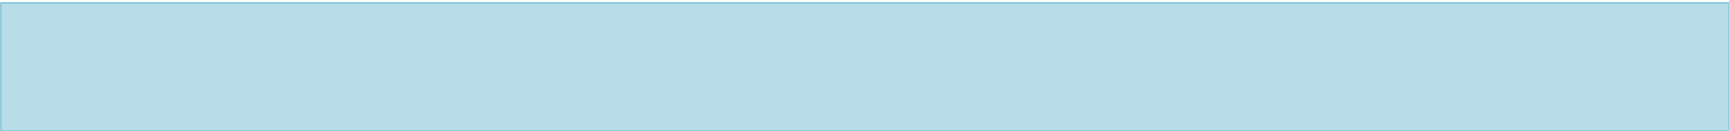
Control Tree:



The control tree depicts the nesting structure of the program.

# Uses of Control-flow Information

- Loop nesting depth plays an important role in optimization heuristics.
  - Deeply nested loops pay off the most for optimization.
- Need to know loop headers / back edges for doing
  - loop invariant code motion
  - loop unrolling
- Dominance information also plays a role in converting to SSA form
  - Used internally by LLVM to do register allocation.



Phi nodes  
Alloc “promotion”  
Register allocation

# REVISITING SSA

# Single Static Assignment (SSA)

- LLVM IR names (via `%uids`) *all* intermediate values computed by the program.
- It makes the order of evaluation explicit.
- Each `%uid` is assigned to only once
  - Contrast with the mutable quadruple form
  - Note that dataflow analyses had these `kill[n]` sets because of updates to variables...
- Naïve implementation of backend: map `%uids` to stack slots
- Better implementation: map `%uids` to registers (as much as possible)
- Question: How do we convert a source program to make maximal use of `%uids`, rather than alloca-created storage?
  - two problems: control flow & location in memory
- Then: How do we convert SSA code to x86, mapping `%uids` to registers?
  - Register allocation.

# Alloca vs. %UID

- Current compilation strategy:

```
int x = 3;  
int y = 0;  
x = x + 1;  
y = x + 2;
```



```
%x = alloca i64  
%y = alloca i64  
store i64* %x, 3  
store i64* %y, 0  
%x1 = load %i64* %x  
%tmp1 = add i64 %x1, 1  
store i64* %x, %tmp1  
%x2 = load %i64* %x  
%tmp2 = add i64 %x2, 2  
store i64* %y, %tmp2
```

- Directly map source variables into %uids?

```
int x = 3;  
int y = 0;  
x = x + 1;  
y = x + 2;
```



```
int x1 = 3;  
int y1 = 0;  
x2 = x1 + 1;  
y2 = x2 + 2;
```



```
%x1 = add i64 3, 0  
%y1 = add i64 0, 0  
%x2 = add i64 %x1, 1  
%y2 = add i64 %x2, 2
```

- Does this always work?



# What about If-then-else?

- How do we translate this into SSA?

```
int y = ...  
int x = ...  
int z = ...  
if (p) {  
    x = y + 1;  
} else {  
    x = y * 2;  
}  
z = x + 3;
```



```
entry:  
    %y1 = ...  
    %x1 = ...  
    %z1 = ...  
    %p = icmp ...  
    br i1 %p, label %then, label %else  
then:  
    %x2 = add i64 %y1, 1  
    br label %merge  
else:  
    %x3 = mult i64 %y1, 2  
merge:  
    %z2 = %add i64 ???, 3
```

- What do we put for ???

# Phi Functions

- Solution:  $\phi$  functions
  - Fictitious operator, used only for analysis
    - implemented by Mov at x86 level
  - Chooses among different versions of a variable based on the path by which control enters the phi node.

$\%uid = \text{phi} \langle \text{ty} \rangle v_1, \langle \text{label}_1 \rangle, \dots, v_n, \langle \text{label}_n \rangle$

```
int y = ...
int x = ...
int z = ...
if (p) {
    x = y + 1;
} else {
    x = y * 2;
}
z = x + 3;
```



```
entry:
    %y1 = ...
    %x1 = ...
    %z1 = ...
    %p = icmp ...
    br i1 %p, label %then, label %else
then:
    %x2 = add i64 %y1, 1
    br label %merge
else:
    %x3 = mult i64 %y1, 2
merge:
    %x4 = phi i64 %x2, %then, %x3, %else
    %z2 = %add i64 %x4, 3
```

# Phi Nodes and Loops

- Importantly, the `%uids` on the right-hand side of a phi node can be defined “later” in the control-flow graph.
  - Means that `%uids` can hold values “around a loop”
  - Scope of `%uids` is defined by *dominance*

```
entry:
    %y1 = ...
    %x1 = ...
    br label %body

body:
    %x2 = phi i64 %x1, %entry, %x3, %body
    %x3 = add i64 %x2, %y1
    %p = icmp slt i64, %x3, 10
    br i1 %p, label %body, label %after

after:
    ...
```

# Alloca Promotion

- Not all source variables can be allocated to registers
  - If the address of the variable is taken (as permitted in C, for example)
  - If the address of the variable “escapes” (by being passed to a function)
- An alloca instruction is called promotable if neither of the two conditions above holds

```
entry:
    %x = alloca i64           // %x cannot be promoted
    %y = call malloc(i64 8)
    %ptr = bitcast i8* %y to i64**
    store i65** %ptr, %x      // store the pointer into the heap
```

```
entry:
    %x = alloca i64           // %x cannot be promoted
    %y = call foo(i64* %x)    // foo may store the pointer into the heap
```

- Happily, most local variables declared in source programs are promotable
  - That means they can be register allocated

# Converting to SSA: Overview

- Start with the ordinary control flow graph that uses allocas
  - Identify “promotable” allocas
- Compute dominator tree information
- Calculate def/use information for each such allocated variable
- Place  $\phi$  functions for each variable at necessary “join points”
- Replace loads/stores to alloc’ed variables with freshly-generated `%uids`
- Eliminate the now unneeded load/store/alloca instructions.

# Phi Placement (Inefficient)

- Less efficient, but easier to understand:
- Place phi nodes "maximally" (i.e. at every node with  $> 1$  predecessors)
- If all values flowing into phi node are the same, then eliminate it:  


```
%x = phi t %y, %pred1 %y %pred2 ... %y %predK  
// code that uses %x  
⇒  
// code with %x replaced by %y
```
- Interleave with other optimizations
  - copy propagation
  - constant propagation
  - etc.

# Phi Placement (Efficient)


- Need to calculate the “Dominance Frontier”
- Node A *strictly dominates* node B if A dominates B and  $A \neq B$ .
  - A dominates A but A does not strictly dominate A
- The *dominance frontier* of a node B is the set of all CFG nodes y such that B dominates a predecessor of y but does not strictly dominate y
  - Write  $DF[B]$  for the dominance frontier of node B.
  - Intuitively, which downstream nodes of B **with multiple predecessors** are not strictly dominated by B.

# Dominance Frontiers

- Example of a dominance frontier calculation results
- $DF[1] = \{1\}$ ,  $DF[2] = \{1,2\}$ ,  $DF[3] = \{2\}$ ,  $DF[4] = \{1\}$ ,  $DF[5] = \{8,0\}$ ,  $DF[6] = \{8\}$ ,  $DF[7] = \{7,0\}$ ,  $DF[8] = \{0\}$ ,  $DF[9] = \{7,0\}$ ,  $DF[0] = \{\}$



Control-flow Graph



Dominator Tree



# Algorithm For Computing DF[n]

- Assume that `doms[n]` stores the dominator tree (so that `doms[n]` is the *immediate dominator* of `n` in the tree)
- Adds each `B` to the DF sets to which it belongs

for all nodes `B`

```
    if  $\#(\text{pred}[B]) \geq 2$                                 // (just an optimization)
        for each  $p \in \text{pred}[B]$  {
            runner :=  $p$                                     // start at the predecessor of  $B$ 
            while (runner  $\neq$  doms[B])                    // walk up the tree adding  $B$ 
                DF[runner] := DF[runner]  $\cup$   $\{B\}$ 
                runner := doms[runner]
        }
```

# Insert $\phi$ at Join Points

- Lift the  $DF[n]$  to a set of nodes  $N$  in the obvious way:  
$$DF[N] = \bigcup_{n \in N} DF[n]$$
- Suppose that a variable  $x$  is defined at a set of nodes  $N$ .

$$DF_0[N] = DF[N]$$

$$DF_{i+1}[N] = DF[DF_i[N] \cup N]$$

Let  $J[N]$  be the *least fixed point* of the sequence:

$$DF_0[N] \subseteq DF_1[N] \subseteq DF_2[N] \subseteq DF_3[N] \subseteq \dots$$

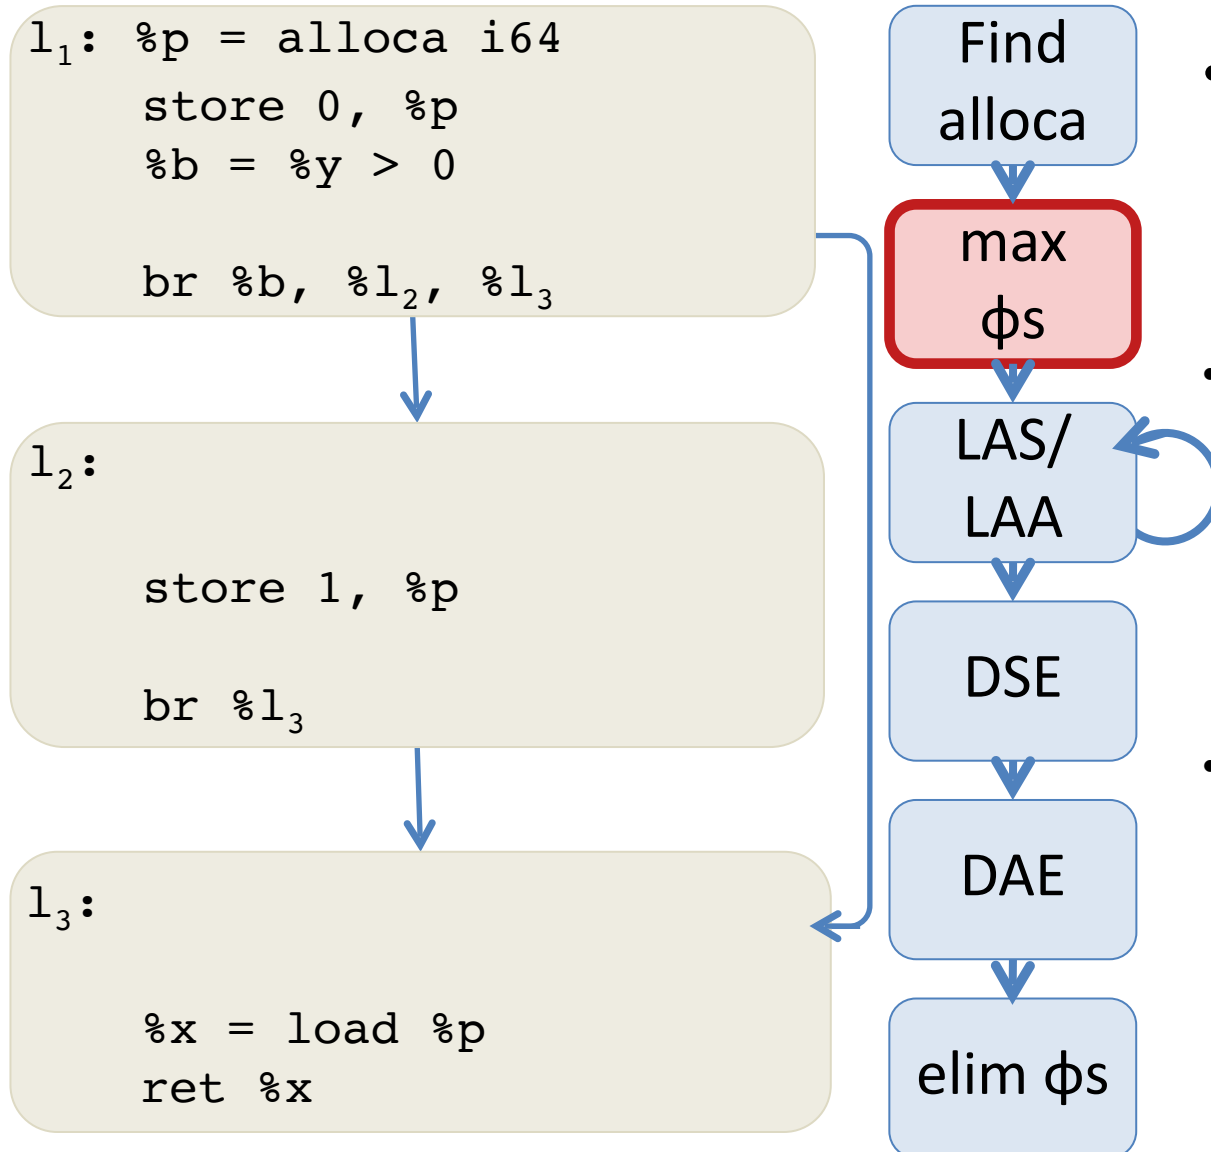
That is,  $J[N] = DF_k[N]$  for some  $k$  such that  $DF_k[N] = DF_{k+1}[N]$

- $J[N]$  is called the “join points” for the set  $N$
- We insert  $\phi$  functions for the variable  $x$  at each node in  $J[N]$ .
  - $x = \phi(x, x, \dots, x)$ ; (one “ $x$ ” argument for each predecessor of the node)
  - In practice,  $J[N]$  is never directly computed, instead you use a worklist algorithm that keeps adding nodes for  $DF_k[N]$  until there are no changes, just as in the dataflow solver.
- Intuition:
  - If  $N$  is the set of places where  $x$  is modified, then  $DF[N]$  is the places where  $\phi$  nodes need to be added, but those also “count” as modifications of  $x$ , so we need to insert the  $\phi$  nodes to capture those modifications too...

# Example Join-point Calculation

- Suppose the variable  $x$  is modified at nodes 3 and 6
  - Where would we need to add phi nodes?
- $DF_0[\{3,6\}] = DF[\{3,6\}] = DF[3] \cup DF[6] = \{2,8\}$
- $DF_1[\{3,6\}]$ 
  - $= DF[DF_0\{3,6\} \cup \{3,6\}]$
  - $= DF[\{2,3,6,8\}]$
  - $= DF[2] \cup DF[3] \cup DF[6] \cup DF[8]$
  - $= \{1,2\} \cup \{2\} \cup \{8\} \cup \{0\} = \{1,2,8,0\}$
- $DF_2[\{3,6\}]$ 
  - $= \dots$
  - $= \{1,2,8,0\}$
- So  $J[\{3,6\}] = \{1,2,8,0\}$  and we need to add phi nodes at those four spots.

# Example SSA Optimizations




- How to place phi nodes without breaking SSA?
- Note: the “real” implementation combines many of these steps into one pass.
  - Places phis directly at the dominance frontier
- This example also illustrates other common optimizations:
  - Load after store/alloca
  - Dead store/alloca elimination

# Example SSA Optimizations

```
l1: %p = alloca i64  
      store 0, %p  
      %b = %y > 0  
      %x1 = load %p  
      br %b, %l2, %l3
```

```
l2:  
      store 1, %p  
      %x2 = load %p  
      br %l3
```

```
l3:  
      %x = load %p  
      ret %x
```




- How to place phi nodes without breaking SSA?
- Insert
  - Loads at the end of each block

# Example SSA Optimizations

$l_1$ : `%p = alloca i64`  
`store 0, %p`  
`%b = %y > 0`  
`%x1 = load %p`  
`br %b, %l2, %l3`

$l_2$ : `%x3 =  $\phi$ [%x1, %l1]`  
`store 1, %p`  
`%x2 = load %p`  
`br %l3`

$l_3$ : `%x4 =  $\phi$ [%x1; %l1, %x2: %l2]`  
`%x = load %p`  
`ret %x`




- How to place phi nodes without breaking SSA?
- Insert
  - Loads at the end of each block
  - Insert  $\phi$ -nodes at each block

# Example SSA Optimizations

```
l1: %p = alloca i64  
      store 0, %p  
      %b = %y > 0  
      %x1 = load %p  
      br %b, %l2, %l3
```

```
l2: %x3 =  $\phi$ [%x1, %l1]  
      store %x3, %p  
      store 1, %p  
      %x2 = load %p  
      br %l3
```

```
l3: %x4 =  $\phi$ [%x1; %l1, %x2: %l2]  
      store %x4, %p  
      %x = load %p  
      ret %x
```



- How to place phi nodes without breaking SSA?
- Insert
  - Loads at the end of each block
  - Insert  $\phi$ -nodes at each block
  - Insert stores after  $\phi$ -nodes

# Example SSA Optimizations

```
l1: %p = alloca i64  
      store 0, %p  
      %b = %y > 0  
      %x1 = load %p  
      br %b, %l2, %l3
```

```
l2: %x3 =  $\phi$ [%x1, %l1]  
      store %x3, %p  
      store 1, %p  
      %x2 = load %p  
      br %l3
```

```
l3: %x4 =  $\phi$ [%x1; %l1, %x2: %l2]  
      store %x4, %p  
      %x = load %p  
      ret %x
```

Find  
alloca

max  $\phi$ s

LAS/  
LAA

DSE


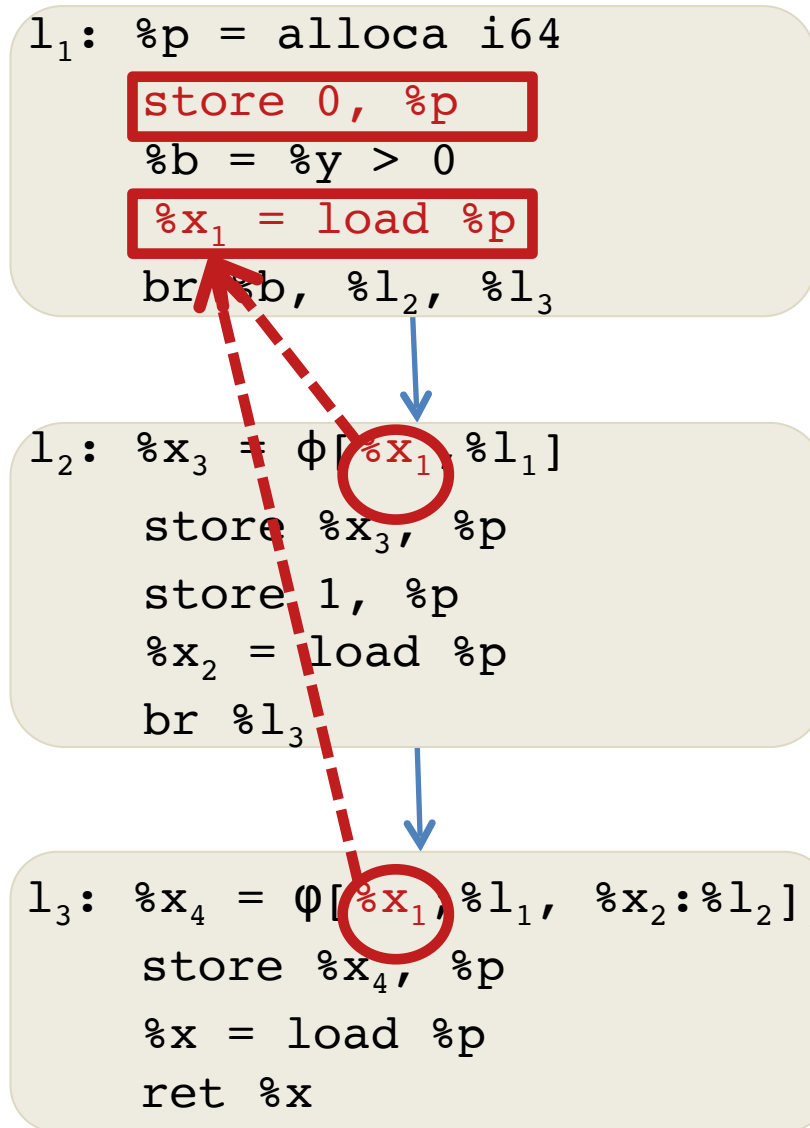
DAE

elim  $\phi$ s

- For loads after stores (LAS):
  - Substitute all uses of the load by the value being stored
  - Remove the load



# Example SSA Optimizations



- For loads after stores (LAS):
  - Substitute all uses of the load by the value being stored
  - Remove the load

# Example SSA Optimizations

$l_1$ : %p = alloca i64

~~store 0, %p~~

%b = %r > 0

~~%x<sub>1</sub> = load %p~~

br %b, %l<sub>2</sub>, %l<sub>3</sub>

$l_2$ : %x<sub>3</sub> =  $\phi[0, \%l_1]$

store %x<sub>3</sub>, %p

store 1, %p

%x<sub>2</sub> = load %p

br %l<sub>3</sub>

$l_3$ : %x<sub>4</sub> =  $\phi[0; \%l_1, \%x_2: \%l_2]$

store %x<sub>4</sub>, %p

%x = load %p

ret %x

Find  
alloca

max  $\phi$ s

LAS/  
LAA

DSE

DAE

elim  $\phi$ s

- For loads after stores (LAS):
  - Substitute all uses of the load by the value being stored
  - Remove the load

# Example SSA Optimizations

```
l1: %p = alloca i64  
      store 0, %p  
      %b = %y > 0  
  
      br %b, %l2, %l3
```

```
l2: %x3 =  $\phi$ [0, %l1]  
      store %x3, %p  
      store 1, %p  
      %x2 = load %p  
      br %l3
```

```
l3: %x4 =  $\phi$ [0; %l1, %x2, %l2]  
      store %x4, %p  
      %x = load %p  
      ret %x
```

Find  
alloca

max  $\phi$ s

LAS/  
LAA

DSE

DAE

elim  $\phi$ s

- For loads after stores (LAS):
  - Substitute all uses of the load by the value being stored
  - Remove the load

# Example SSA Optimizations

```
l1: %p = alloca i64  
      store 0, %p  
      %b = %y > 0  
  
      br %b, %l2, %l3
```

```
l2: %x3 =  $\phi$ [0, %l1]  
      store %x3, %p  
      store 1, %p  
      %x2 = load %p  
      br %l3
```

```
l3: %x4 =  $\phi$ [0; %l1, 1: %l2]  
      store %x4, %p  
      %x = load %p  
      ret %x
```

Find  
alloca

max  $\phi$ s

LAS/  
LAA

DSE

DAE

elim  $\phi$ s

- For loads after stores (LAS):
  - Substitute all uses of the load by the value being stored
  - Remove the load

# Example SSA Optimizations

```
l1: %p = alloca i64  
      store 0, %p  
      %b = %y > 0  
  
      br %b, %l2, %l3
```

```
l2: %x3 =  $\phi$ [0,%l1]  
      store %x3, %p  
      store 1, %p  
  
      br %l3
```

```
l3: %x4 =  $\phi$ [0;%l1, 1;%l2]  
      store %x4, %p  
      %x = load %p  
      ret %x
```

Find  
alloca

max  $\phi$ s

LAS/  
LAA

DSE

DAE

elim  $\phi$ s

- For loads after stores (LAS):
  - Substitute all uses of the load by the value being stored
  - Remove the load

# Example SSA Optimizations

```
l1: %p = alloca i64  
      store 0, %p  
      %b = %y > 0  
  
      br %b, %l2, %l3
```

```
l2: %x3 =  $\phi$ [0,%l1]  
      store %x3, %p  
      store 1, %p  
  
      br %l3
```

```
l3: %x4 =  $\phi$ [0;%l1, 1:%l2]  
      store %x4, %p  
      %x = load %p  
      ret %x4
```

Find  
alloca

max  $\phi$ s

LAS/  
LAA

DSE

DAE

elim  $\phi$ s

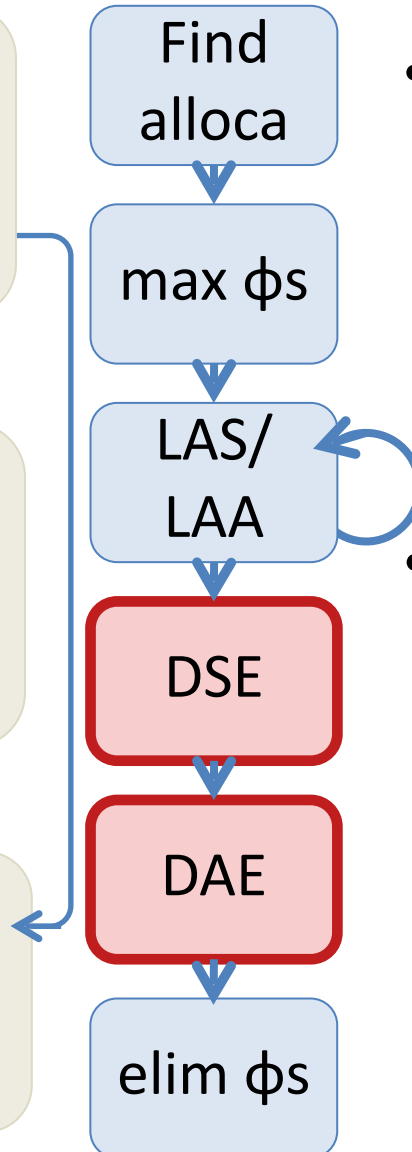
- For loads after stores (LAS):
  - Substitute all uses of the load by the value being stored
  - Remove the load

# Example SSA Optimizations

```
l1: %p = alloca i64  
      store 0, %p  
      %b = %y > 0  
  
      br %b, %l2, %l3
```


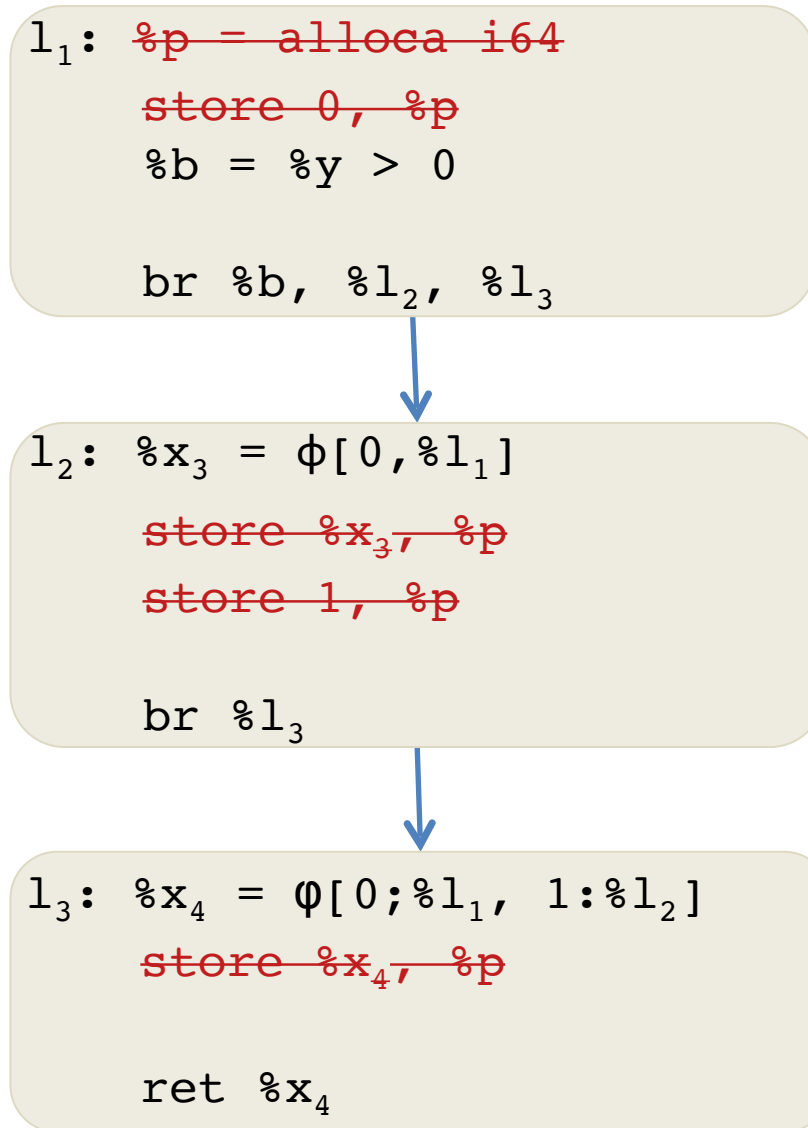
```
l2: %x3 =  $\phi[0, \%l_1]$   
      store %x3, %p  
      store 1, %p  
  
      br %l3
```

```
l3: %x4 =  $\phi[0; \%l_1, 1: \%l_2]$   
      store %x4, %p  
  
      ret %x4
```



- Dead Store Elimination (DSE)
  - Eliminate all stores with no subsequent loads.
- Dead Alloca Elimination (DAE)
  - Eliminate all allocas with no subsequent loads/stores.


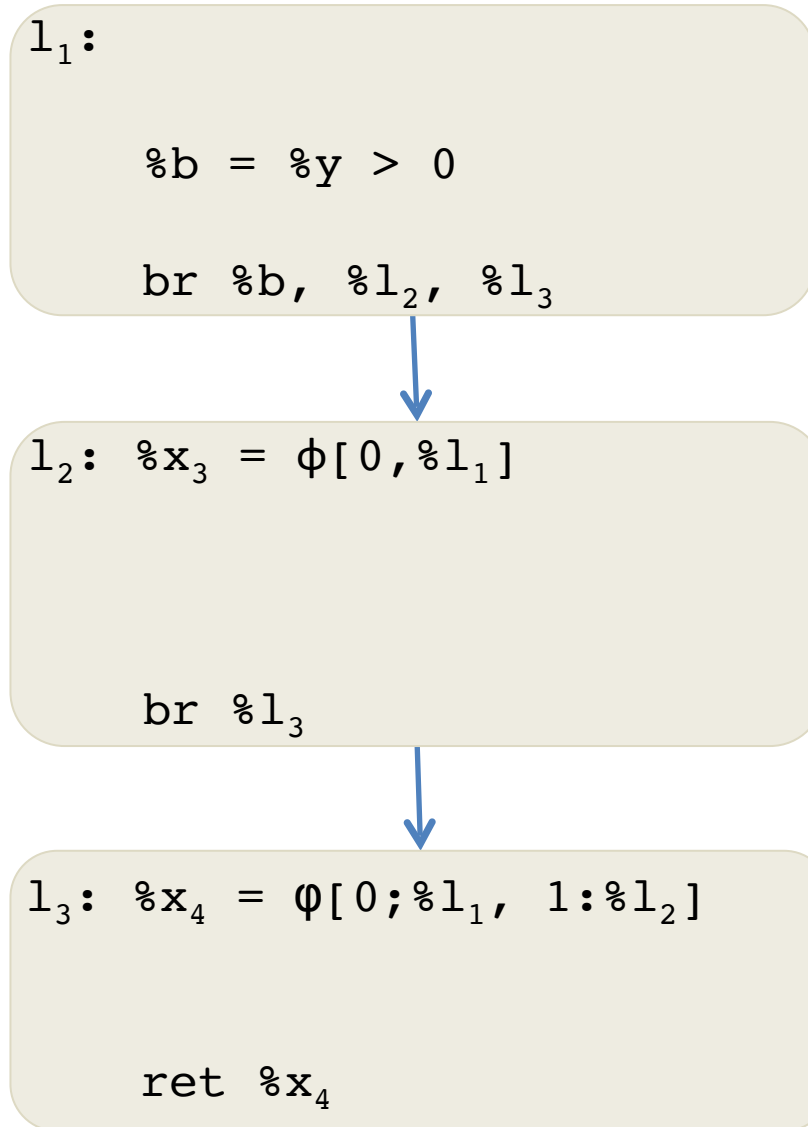
# Example SSA Optimizations



- Dead Store Elimination (DSE)
  - Eliminate all stores with no subsequent loads.
- Dead Alloca Elimination (DAE)
  - Eliminate all allocas with no subsequent loads/stores.


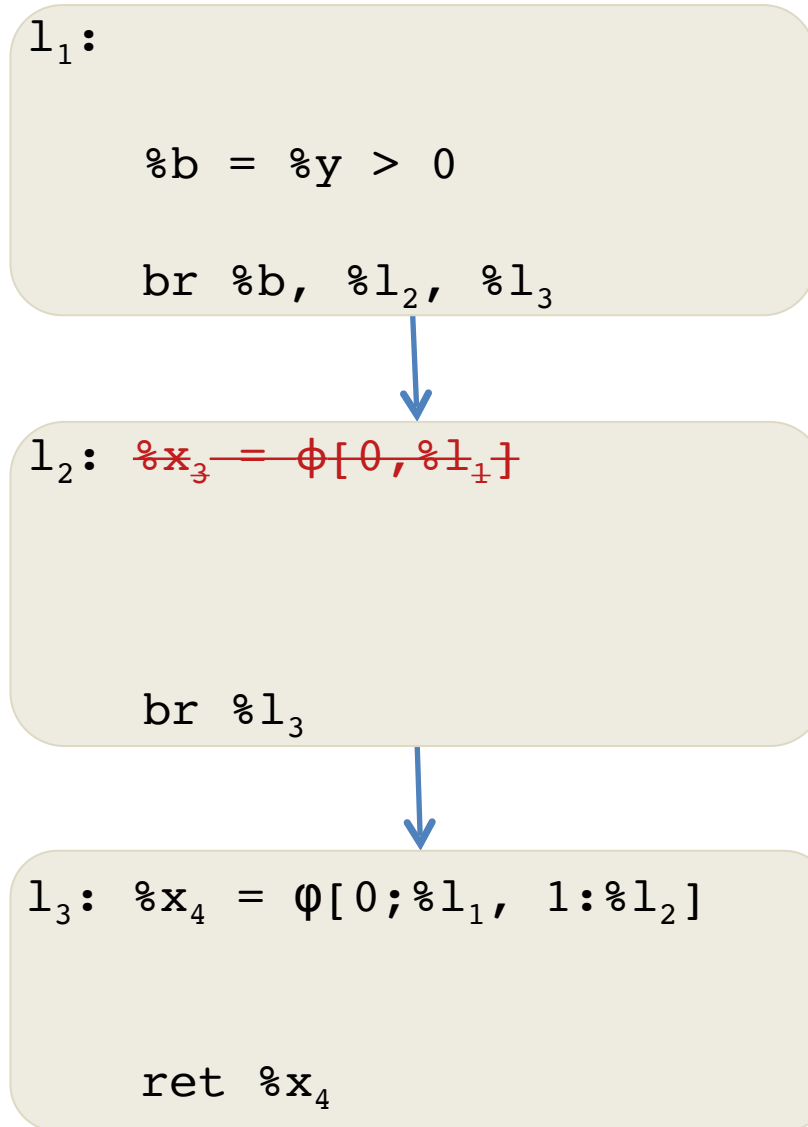


# Example SSA Optimizations



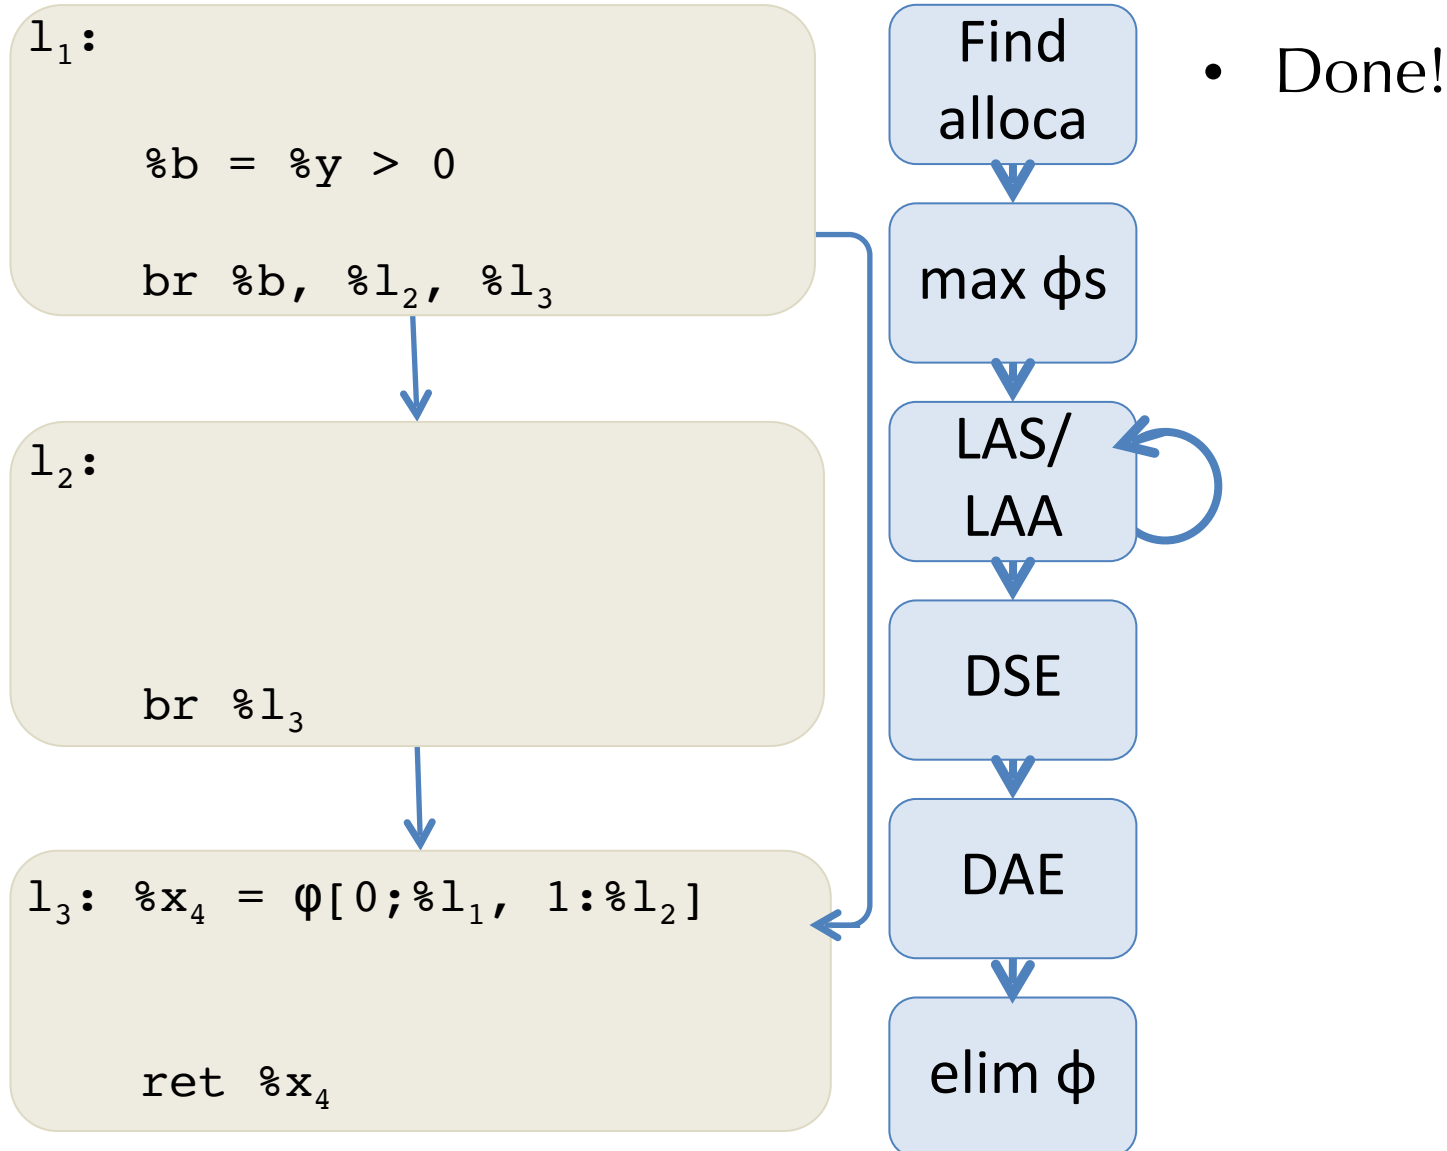
- Eliminate  $\phi$  nodes:
  - Singletons
  - With identical values from each predecessor
  - See Aycock & Horspool, 2002

# Example SSA Optimizations



- Eliminate  $\phi$  nodes:
  - Singletons
  - With identical values from each predecessor

# Example SSA Optimizations



# LLVM Phi Placement

- This transformation is also sometimes called register promotion
  - older versions of LLVM called this “mem2reg” memory to register promotion
- In practice, LLVM combines this transformation with *scalar replacement of aggregates* (SROA)
  - i.e. transforming loads/stores of structured data into loads/stores on register-sized data
- These algorithms are (one reason) why LLVM IR allows annotation of predecessor information in the .ll files
  - Simplifies computing the DF